

# DD2424 Assignment 1 basic

Jacob Malmberg, 198508060558

March 2019

## 1 Analytical gradient check

The analytical gradients were checked against numerical gradients computed by ComputeGradsNumSlow.m. This was done for a batch containing the first 100 datapoints of data\_batch\_1.mat. All 3072 dimensions were included. The relative error was computed following the guidelines laid out on page 7 of Assignment 1:

$$\frac{|g_a - g_n|}{\max(\epsilon, |g_a| + |g_n|)} \quad (1)$$

The results are summarized in 1. According to <http://cs231n.github.io/neural-networks-3/#gradcheck> these are good values which makes me conclude that my gradient computations are bug free. Adding to this, I implemented backward/forward pass using efficient matrix computations as specified in the lecture ppt's.

Configuration	$\lambda$	Max W relative error	Max B relative error
100 datapts	0	5.17e-6	6.12e-8
100 datapts	0.1	9.11 e-6	2.31e-7

Table 1: Table showing summary of relative errors

## 2 Results

The following sections contain the results. The last subsection contains a comment on the effect of regularization and the importance of the correct learning rate.

### 2.1 $\lambda = 0$ , n\_epochs = 40, n\_batch = 100, eta = .1

The accuracy was 27.07%. Loss and cost for this configuration can be found in fig 1. Weights visualization can be found in fig 2.

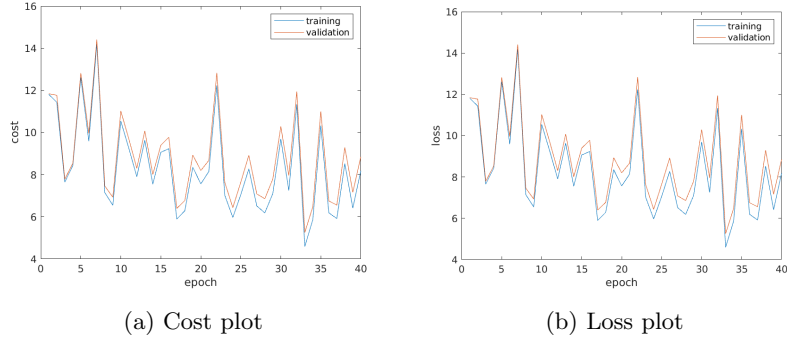


Figure 1: Loss and cost plots for  $\lambda = 0$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = 0.1$



Figure 2: Learnt  $W$  matrix for  $\lambda = 0$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = .1$

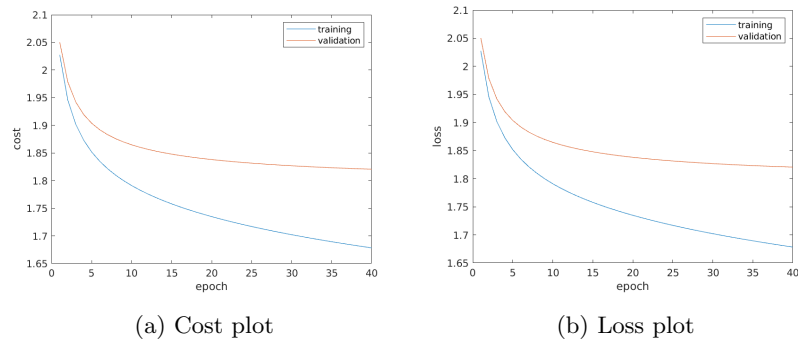


Figure 3: Loss and cost plots for  $\lambda = 0$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = .01$




Figure 4: Learnt W matrix for  $\lambda = 0$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = .01$

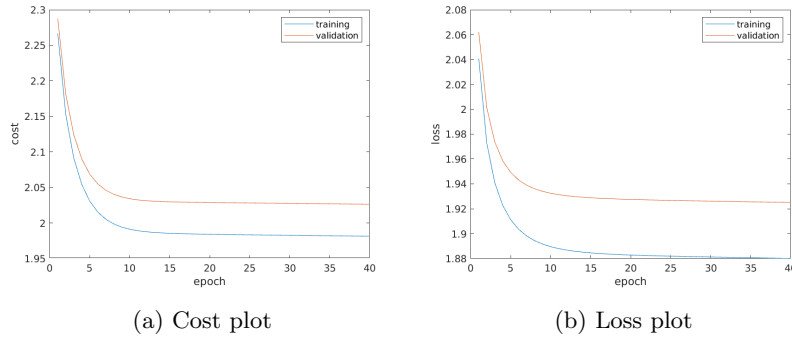


Figure 5: Loss and cost plots for  $\lambda = 0.1$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = 0.01$

## 2.2 $\lambda = 0$ , $n\_epochs = 40$ , $n\_batch = 100$ , $\eta = .01$

The accuracy was 36.65%. Loss and cost for this configuration can be found in fig 3. Weights visualization can be found in fig 4.

## 2.3 $\lambda = 0.1$ , $n\_epochs = 40$ , $n\_batch = 100$ , $\eta = .01$

The accuracy was 33.37%. Loss and cost for this configuration can be found in fig 5. Weights visualization can be found in fig 6.




Figure 6: Learnt W matrix for  $\lambda = 0.1$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = 0.01$

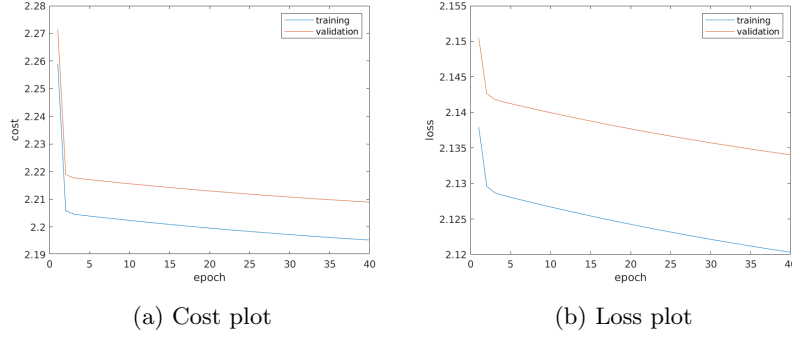


Figure 7: Loss and cost plots for  $\lambda = 1$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = 0.01$



Figure 8: L learnt  $W$  matrix for  $\lambda = 1$ ,  $n\_epochs = 40$ ,  $n\_batch = 100$ ,  $\eta = .01$

## 2.4 $\lambda = 1$ , $n\_epochs = 40$ , $n\_batch = 100$ , $\eta = 0.01$

The accuracy was 21.92%. Loss and cost for this configuration can be found in fig 7. Weights visualization can be found in fig 8.

## 2.5 Effects of regularization and correct learning rate

If the learning rate is too high, we may end up overshooting the minimum point and instead start to climb toward a higher function value when doing gradient descent. I believe this is what is happening in fig. 1. A lower learning rate is slower but should take us toward the minimum. Thus, the learning rate should be well calibrated.

Increasing regularization should prevent overfitting to the training dataset, however a too high  $\lambda$  may result in poor test performance. In fig. 7 the cost for training and validation set is closer to each other than in plots with lower regularization (the y-axis are different in the plots which makes it a little funny looking). With no regularization we may end up with clear overfitting, shown by training cost and validation loss diverging. This is shown in fig 3.